

Lecture 11: Fenchel Conjugate, Dual formulation of the Empirical Risk Minimization, and SDCA

1 Fenchel Conjugate

Definition 1 (Fenchel Conjugate). *Consider a function $f(\cdot)$, then the Fenchel Conjugate is defined to be*

$$f^*(y) = \sup_{x \in \text{dom}(f)} (y^\top x - f(x)).$$

Claim. *The conjugate function $f^*(y)$ is always convex, even if $f(\cdot)$ is non-convex.*

Proof. Let $h_x(y) := y^\top x - f(x)$. Observe that h_x is an affine function of y and therefore also convex. Let $\alpha \in [0, 1]$ and $y_1, y_2 \in \text{dom}(f^*)$. Then, we have

$$\begin{aligned} f^*((1-\alpha)y_1 + \alpha y_2) &= \sup_{x \in \text{dom}(f)} h_x((1-\alpha)y_1 + \alpha y_2) \\ &= \sup_{x \in \text{dom}(f)} (1-\alpha)h_x(y_1) + \alpha h_x(y_2) \\ &\leq (1-\alpha) \sup_{x \in \text{dom}(f)} h_x(y_1) + \alpha \sup_{x \in \text{dom}(f)} h_x(y_2) \\ &= (1-\alpha)f^*(y_1) + \alpha f^*(y_2). \end{aligned}$$

Thus, by the zero-order characterization of convexity, we have that f^* is convex. \square

Exercise 1. $f(x) = a^\top x + b$.

$$\begin{aligned} f^*(y) &= \sup_{x \in \text{dom}(f)} (\langle y, x \rangle - f(x)) \\ &= \sup_{x \in \text{dom}(f)} (\langle y - a, x \rangle - b) \\ &= \begin{cases} -b & , \text{if } y = a \\ \infty & , \text{otherwise} \end{cases}. \end{aligned}$$

Exercise 2. $f(x) = \frac{1}{2}x^2$.

$$\begin{aligned}
f^*(y) &= \sup_{x \in \text{dom}(f)} (\langle y, x \rangle - f(x)) \\
&= \sup_{x \in \text{dom}(f)} \left(\langle y, x \rangle - \frac{1}{2}x^2 \right)
\end{aligned}$$

Let $h(x) := \langle y, x \rangle - \frac{1}{2}x^2$. Then, the maximizer of h can be found as

$$\nabla h(x) = 0 \Leftrightarrow x = y$$

Thus,

$$f^*(y) = y^2 - \frac{1}{2}y^2 = \frac{1}{2}y^2$$

1.1 Fenchel inequality

By the definition of the conjugate function, we have the following result:

Theorem 1 (Fenchel inequality). *For any x and y , we have*

$$f^*(y) \geq y^\top x - f(x).$$

Question. When do we have the equality?

$$f^*(y) + f(x) = y^\top x$$

In the following, we are going to answer this equation and prove the following theorem:

Theorem 2. *If $f(\cdot)$ is closed and convex, then the following are equivalent:*

- i. $f^*(y) + f(x) = y^\top x$.
- ii. $y = \nabla f(x)$.
- iii. $x = \nabla f^*(y)$.

Now, recall the definition of open and closed sets.

Definition 2 (Open set). *A set S is open if it contains an open ball about each of its points. That is, for all $x \in S$, there exists $\epsilon > 0$ such that $B_\epsilon(x) \subset S$.*

Definition 3 (Closed set). *A set S is closed if its complement is open.*

We will now introduce the definition of closed functions.

Definition 4 (Closed function). *A function is closed if its sublevel set is a closed set, i.e.,*

$$\{x \in \text{dom}(f) : f(x) \leq \alpha\}$$

is a closed set.

Counterexample. $f(x) = \exp(-x)$ is not a closed function. Observe that its sublevel set $\{x \in \text{dom}(f) : \exp(-x) \leq \alpha\}$ is not closed.

1.2 The inverse of the gradient map

Theorem 3. *Suppose that $f(\cdot)$ is closed and convex. Then, $y = \nabla f(x)$ if and only if $x = \nabla f^*(y)$.*

Proof. We will only prove the " \Rightarrow " direction, that is we will show that if $y = \nabla f(x)$ then $x = \nabla f^*(y)$. Let $y = \nabla f(x)$. By the first-order characterization of convexity, for any $u \in \mathbb{R}^d$ we have

$$f(u) \geq f(x) + \langle y, u - x \rangle.$$

Additionally, we have

$$\begin{aligned} f^*(y) &= \sup_u (\langle u, y \rangle - f(u)) && \text{(by definition of conjugate function)} \\ &\leq \sup_u \langle u, y \rangle - (f(x) + \langle y, u - x \rangle) && \text{(by convexity)} \\ &= \langle x, y \rangle - f(x) \end{aligned} \tag{1}$$

Recall that for a convex function $h(\cdot)$ defined over a convex set C , a vector g_x is said to be a sub-gradient of $h(\cdot)$ at a point $x \in C$ if for any $y \in C$

$$h(y) \geq h(x) + \langle g_x, y - x \rangle.$$

Now, for any $z \in \mathbb{R}^d$ we have

$$\begin{aligned} f^*(z) &\geq \langle z, x \rangle - f(x) && \text{(by definition of the Fenchel inequality)} \\ &= \langle z - y, x \rangle - f(x) + \langle y, x \rangle \\ &\geq \langle z - y, x \rangle + f^*(y) && \text{(by inequality (1))} \end{aligned}$$

By the fact that $f^*(\cdot)$ is convex (and differentiable) and by the definition of the subgradient we have that

$$x = \nabla f^*(y),$$

which concludes the proof.

To prove the other direction, we follow the same lines of the proof as above. Specifically, we let $r := f^*$ and the function $r(\cdot)$ is convex and closed. So we can use the above argument for $r(\cdot)$ and deduce that if $x = \nabla r(y)$, then $y = \nabla r^*(x)$. Then, using the fact that if $f(\cdot)$ is closed and convex, then the bi-conjugate $f^{**}(x)$ is equal to the original function $f(\cdot)$ itself, we can complete the proof

□

Theorem 4. *If $f(\cdot)$ is closed and convex, then the following are equivalent:*

$$f^*(y) + f(x) = y^\top x \iff y = \nabla f(x) \iff x = \nabla f^*(y).$$

Proof. We have shown the equivalency of the last two items in Theorem 3.

To show the equivalency of the first two items, we would like to figure out the answer to the following questions:

Question 1. What is $\arg \sup_{x \in \text{dom}(f)} (y^\top x - f(x))$?

This is because the $\arg \sup_{x \in \text{dom}(f)} (y^\top x - f(x))$ is what makes the Fenchel inequality becomes the equality.

As the function $f(\cdot)$ is closed, we have that the $\arg \sup$ becomes $\arg \max$. Therefore,

$$\arg \sup_{x \in \text{dom}(f)} \underbrace{(y^\top x - f(x))}_{q(x)} = \arg \max_x (y^\top x - f(x)) = \nabla f^*(y),$$

because

$$\nabla q(x) = 0 \iff y = \nabla f(x) \iff x = \nabla f^*(y).$$

The last step follows from Theorem 3. We can conclude that $y = \nabla f(x) \iff x = \nabla f^*(y)$ if and only if the Fenchel inequality becomes the equality, when $f(\cdot)$ is closed and convex.

Question 2. What is $\arg \sup_{y \in \text{dom}(f^*)} (y^\top x - f^*(y))$?

$$\arg \max_y \underbrace{(y^\top x - f^*(y))}_{s(y)} = \nabla f(x),$$

because

$$\nabla s(y) = 0 \Leftrightarrow x = \nabla f^*(y) \Leftrightarrow y = \nabla f(x).$$

The last step follows from Theorem 3.

□

2 Regularized Empirical Risk Minimization

If the primal problem is

$$\min_{x \in \mathbb{R}^d} F(x), \quad \text{where } F(x) := \frac{1}{n} \sum_{i=1}^n f_i(x^\top z_i) + \frac{\lambda}{2} \|x\|_2^2,$$

then the dual problem is

$$\max_{\alpha \in \mathbb{R}^n} D(\alpha), \quad \text{where } D(\alpha) := \frac{1}{n} \sum_{i=1}^n -f_i^*(-\alpha_i) - \frac{\lambda}{2} \left\| \frac{1}{\lambda n} \sum_{i=1}^n \alpha_i z_i \right\|_2^2.$$

We will show how the dual problem is derived from the primal problem. Consider the following constrained optimization problem

$$\begin{aligned} \min_{x \in \mathbb{R}^d} \sum_{i=1}^n f_i(\theta_i) + \frac{\lambda n}{2} \|x\|_2^2 \\ \text{subject to } \forall i, \theta_i = z_i^\top x, \end{aligned}$$

where we have introduced variables $\{\theta_i\}_{i=1}^n$.

Step 1. Constructing the Lagrangian

The Lagrangian is formulated as

$$L(x, \{\theta_i\}, \{\alpha_i\}) = \sum_{i=1}^n \left[f_i(\theta_i) + \alpha_i (\theta_i - z_i^\top x) \right] + \frac{\lambda n}{2} \|x\|_2^2$$

Step 2. Optimizing over primal variables to get the dual function

We have that

$$\begin{aligned} \min_{x, \theta_1, \dots, \theta_n} \sum_{i=1}^n \left(f_i(\theta_i) + \alpha_i \theta_i - \alpha_i z_i^\top x \right) + \frac{\lambda n}{2} \|x\|_2^2 \\ \iff \min_x \sum_{i=1}^n \left(\min_{\theta_i} f_i(\theta_i) + \alpha_i \theta_i \right) + \frac{\lambda n}{2} \|x\|_2^2 - \sum_{i=1}^n \alpha_i z_i^\top x. \end{aligned}$$

Now, observe that

$$\min_{\theta} q(\theta) = - \max_{\theta} (-q(\theta)) .$$

Thus, we have that

$$\begin{aligned} \left(\min_{\theta_i} f_i(\theta_i) + \alpha_i \theta_i \right) &= - \max_{\theta_i} \left[- (f_i(\theta_i) + \alpha_i \theta_i) \right] \\ &= - \max_{\theta_i} [-\alpha_i \theta_i - f_i(\theta_i)] \\ &= -f_i^*(\alpha_i) \end{aligned} \quad \text{(by definition of the conjugate)}$$

Therefore, using the above result we can rewrite

$$\begin{aligned} \min_{x, \theta_1 - \theta_n} \sum_{i=1}^n \left(f_i(\theta_i) + \alpha_i \theta_i - \alpha_i z_i^\top x \right) + \frac{\lambda n}{2} \|x\|_2^2 \\ \iff \min_x \sum_{i=1}^n \left(\min_{\theta_i} f_i(\theta_i) + \alpha_i \theta_i \right) + \frac{\lambda n}{2} \|x\|_2^2 - \sum_{i=1}^n \alpha_i z_i^\top x \\ \iff - \sum_{i=1}^n f_i^*(-\alpha_i) + \underbrace{\min_x \frac{\lambda n}{2} \|x\|_2^2 - \sum_{i=1}^n \alpha_i z_i^\top x}_{q(x)} . \end{aligned}$$

Additionally, observe that

$$q(x) = 0 \iff \lambda n x = \sum_{i=1}^n \alpha_i z_i \iff x = \frac{1}{\lambda n} \sum_{i=1}^n \alpha_i z_i$$

The equation

$$x = \frac{1}{\lambda n} \sum_{i=1}^n \alpha_i z_i$$

describes the **relation between primal variables and dual variables**. Using this

result we have

$$\begin{aligned}
\min_x \frac{\lambda n}{2} \|x\|_2^2 - \sum_{i=1}^n \alpha_i z_i^\top x &= \frac{\lambda n}{2} \left\| \frac{1}{\lambda n} \sum_{i=1}^n \alpha_i z_i \right\|_2^2 - \left\langle \sum_{i=1}^n \alpha_i z_i, \frac{1}{\lambda n} \sum_{i=1}^n \alpha_i z_i \right\rangle \\
&= \frac{1}{2\lambda n} \left\| \sum_{i=1}^n \alpha_i z_i \right\|_2^2 - \frac{1}{\lambda n} \left\| \sum_{i=1}^n \alpha_i z_i \right\|_2^2 \\
&= -\frac{1}{2\lambda n} \left\| \sum_{i=1}^n \alpha_i z_i \right\|_2^2 \\
&= -\frac{\lambda n}{2} \left\| \frac{1}{\lambda n} \sum_{i=1}^n \alpha_i z_i \right\|_2^2.
\end{aligned}$$

Plugging this in the objective we get

$$\begin{aligned}
\min_{x, \theta_1 - \theta_n} \sum_{i=1}^n \left(f_i(\theta_i) + \alpha_i \theta_i - \alpha_i z_i^\top x \right) + \frac{\lambda n}{2} \|x\|_2^2 \\
\iff \min_x \sum_{i=1}^n \left(\min_{\theta_i} f_i(\theta_i) + \alpha_i \theta_i \right) + \frac{\lambda n}{2} \|x\|_2^2 - \sum_{i=1}^n \alpha_i z_i^\top x \\
\iff -\sum_{i=1}^n f_i^*(-\alpha_i) + \min_x \frac{\lambda n}{2} \|x\|_2^2 - \sum_{i=1}^n \alpha_i z_i^\top x \\
\iff \underbrace{-\sum_{i=1}^n f_i^*(-\alpha_i) - \frac{\lambda n}{2} \left\| \frac{1}{\lambda n} \sum_{i=1}^n \alpha_i z_i \right\|_2^2}_{D(\alpha)}.
\end{aligned}$$

Step 3. Solve $\max_{\alpha \in \mathbb{R}^n} D(\alpha)$

3 Duality Gap

Recall from the previous section that the relation between primal variables and dual variables is

$$x = \frac{1}{\lambda n} \sum_{i=1}^n \alpha_i z_i.$$

The duality gap is defined by

$$\text{Duality gap} := F(x(\alpha)) - D(\alpha)$$

Then, the primal optimality gap $F(x(\alpha)) - F_*$ is bounded by the duality gap $:= F(x(\alpha)) - D(\alpha)$.

Remark: This reveals the benefit of considering developing algorithms in the dual space. Since we can obtain an upper-bound of the optimality gap on the fly during the execution of the underlying dual algorithm. We demonstrate one of the classical algorithms in the next section.

4 Stochastic Dual Coordinate Ascent (SDCA)

4.1 Main Idea

Consider the unconstrained optimization problem we introduced

$$\max_{\alpha \in \mathbb{R}^n} D(\alpha), \quad \text{where } D(\alpha) := \frac{1}{n} \sum_{i=1}^n -f_i^*(-\alpha_i) - \frac{\lambda}{2} \left\| \frac{1}{\lambda n} \sum_{i=1}^n \alpha_i z_i \right\|_2^2.$$

Consider updating a dual variable $\alpha_i \in \mathbb{R}^n$ at a time. That is, at the k -th iteration, we pick $i_k \in [n]$. Then, we have

$$\begin{aligned} & \max_{\alpha_{i_k}} -\frac{1}{n} f_{i_k}^*(-\alpha_{i_k}) - \frac{\lambda}{2} \left\| \frac{1}{\lambda n} \sum_{i=1}^n \alpha_i z_i \right\|_2^2 \\ \iff & \max_{\alpha_{i_k}} -\frac{1}{n} f_{i_k}^*(-\alpha_{i_k}) - \frac{\lambda}{2} \left\| \frac{1}{\lambda n} \sum_{i=1}^n \alpha_i^{(k-1)} z_i + \frac{1}{\lambda n} \Delta \alpha_{i_k} z_{i_k} \right\|_2^2 \\ \iff & \max_{\Delta \alpha_{i_k}} -\frac{1}{n} f_{i_k}^* \left(- \left(\alpha_{i_k}^{(k-1)} + \Delta \alpha_{i_k} \right) \right) - \frac{\lambda}{2} \left\| x^{(k-1)} + \frac{1}{\lambda n} \Delta \alpha_{i_k} z_{i_k} \right\|_2^2, \end{aligned}$$

where

$$\alpha_{i_k} = \underbrace{\alpha_{i_k}^{(k-1)}}_{\text{fixed}} + \underbrace{\Delta \alpha_{i_k}}_{\text{variable}}$$

and

$$x^{(k-1)} = \frac{1}{\lambda n} \sum_{i=1}^n \alpha_i^{(k-1)} z_i.$$

4.2 Algorithm

Below is a formal statement of the SDCA algorithm [3].

Algorithm 1 Stochastic Dual Coordinate Ascent (SDCA)

- 1: Init dual variables $\alpha^{(1)} \in \mathbb{R}^n$.
- 2: **for** $k = 1, 2, \dots, K$ **do**
- 3: Randomly pick a dual coordinate $i_k \in [n]$.
- 4: Maximizes the dual problem by updating the dual variable i_k while fixing the others

$$\max_{\Delta\alpha_{i_k}} -\frac{1}{n}f_{i_k}^* \left(- \left(\alpha_{i_k}^{(k-1)} + \Delta\alpha_{i_k} \right) \right) - \frac{\lambda}{2} \left\| x^{(k-1)} + \frac{1}{\lambda n} \Delta\alpha_{i_k} z_{i_k} \right\|_2^2.$$

- 5: $\alpha^{(k)} = \alpha^{(k-1)} + \Delta\alpha_{i_k} e_{i_k} \in \mathbb{R}^n$.
 - 6: $x^{(k)} = x^{(k-1)} + \frac{1}{\lambda n} \Delta\alpha_{i_k} z_{i_k} \in \mathbb{R}^d$.
 - 7: **end for**
 - 8: Output: $x(\alpha^{(K)}) := \frac{1}{\lambda n} \sum_{i=1}^n \alpha_i^{(K)} z_i$.
-

Remark: Note that in the primal space, each primal coordinate corresponds to a dimension of the “feature” vector; on the other hand, in the dual space, a dual coordinate corresponds to a data point. Randomly picking up a dual coordinate to update is about randomly choosing a sample to use for the update.

4.3 Example

Example Let us consider $f_i(\theta) := \max\{1 - y_i\theta\}$ being the hinge loss, where $y_i \in \{-1, +1\}$. Its conjugate function is

$$f_i^*(a) = \begin{cases} ay_i & , \text{ if } ay_i \in [-1, 0], \\ \infty & , \text{ otherwise} \end{cases}.$$

The update of the SDCA for the hinge loss is

$$\Delta\alpha_{i_k} = y_{i_k} \max \left(0, \max \left(1, \frac{1 - z_{i_k}^\top x^{(k-1)} y_{i_k}}{\|z_{i_k}\|_2^2 / \lambda n} + \alpha_{i_k}^{(k-1)} y_{i_k} \right) \right) - \alpha_{i_k}^{(k-1)}.$$

Bibliographic notes

For references on conjugate functions, please refer to Chapter 5 of Algorithms for Convex Optimization by Nisheeth K. Vishnoi [1] and Chapter 5 of Convex Optimization by Stephen Boyd and Lieven Vandenbergh.

References

- [1] Nisheeth K. Vishnoi. Algorithms for Convex Optimization. Cambridge University Press, 2021.
- [2] Stephen Boyd and Lieven Vandenberghe, Convex Optimization Cambridge University Press, 2004.
- [3] Rie Johnson and Tong Zhang Accelerating Stochastic Gradient Descent using Predictive Variance Reduction. NeurIPS 2013.